

Govern.ai

The operating layer for trustworthy HR work

Specialized agents act. Policy evidence grounds.
Human checkpoints keep sensitive decisions in control.

Powered by Fireworks AI and AMD



HR automation fails when action outruns evidence

A chatbot can answer. A real operating layer must coordinate work, show why it acted, and know when to stop.

“Why did this happen?”

In HR, the answer must be more than
‘the model said so.’

The product thesis

Fleet ♦ evidence ♦ governance

01

Action without context

A model can trigger a workflow without showing the policy or case evidence behind it.

02

Automation without a stop

Sensitive actions need a human checkpoint—not a hidden confidence threshold.

03

Inference without a record

If the decision is not captured with who, when, and why, it cannot be audited.

Three legs turn ‘an AI’ into an HR operating layer

Remove any one of them and the system collapses into a toy, a search box, or a black box.




Govern.ai = action that can explain itself and stop safely

One console coordinates five workloads—with a human boundary

The system is designed for HR operators, not as a general-purpose employee dashboard.

- **Ticket triage**
Routes routine and urgent work; urgent cases escalate.
- **Policy Q&A**
Answers from versioned evidence; weak evidence triggers review.
- **Resume screening**
Transparent advisory rubric; never auto-rejects.
- **Onboarding**
Runs a checklist and pauses sensitive steps for approval.
- **Attrition advisory**
Surfaces risk factors; never a punitive verdict.



HR OPERATORS

Employees get a focused, chat-first surface; the command center is reserved for the people who operate it.

The proof is an auditable resume review

Decision support is useful only when the handoff is transparent, PII-safe, and human-approved.

01 Intake

Resume arrives as PDF, image, or text.

02 Blind ♦ structure


Identity is blinded; skills and experience become a typed profile.

03 Guard ♦ explain

Policy guard checks the rubric and generates reasons—not a verdict.

04 Approve ♦ audit

A human checkpoint closes the loop and records the decision.



A2A-shaped handoffs carry certified, PII-safe artifacts—not raw text.

Provider routing is explicit, cost-bounded, and reversible

The durable control plane keeps workflows useful even when model inference is unavailable.

Deterministic first

Obvious triage, local scoring, and templated explanations use zero model spend.

Fireworks AI online


Interactive ambiguity, grounded synthesis, and streaming are opt-in and allowlisted.

AMD / vLLM optional

A self-hosted route exists; live hardware claims stay gated by named runtime evidence.


Technical architecture

One provider-neutral control plane; sensitive paths stay behind the application boundary.



Governance is product behavior—not a footer

The system makes the safe path visible: role boundaries, human approvals, redaction, and evidence all live in the workflow.



- **Human checkpoints**
Sensitive actions pause for explicit approval.
- **Least privilege**
Viewer → analyst → manager → admin scopes the surface and the action.
- **PII boundaries**
Audit and provider payloads use redacted representations.
- **Append-only evidence**
Every agent action and human decision is exportable for review.

A deployable reference today; measured depth next

Govern.ai is designed to make trustworthy AI in HR inspectable now—and to earn stronger claims with evidence, not optimism.

Proven in the repository

- Zero-secret deterministic evaluation
- Policy RAG with pgvector + citations
- Human approval gates + immutable audit
- Provider-neutral routing + batch lifecycle

Gates still required

- Credentialed Fireworks AI smoke
- Named AMD/ROCm/vLLM runtime evidence
- Hardware benchmark + numerical error evidence
- Tenant isolation + enterprise identity

Build HR automation people can question, review, and trust.